

Logistic Regression

Use the dataset, perform necessary pre-processing and build a logistic regression model. divide the train data itself into 70-30 ratio and print the performance metrics

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
df = pd.read_csv("/content/telecom_customer_churn.csv")
```

df

	Customer ID	Gender	Age	Married	Number of Dependents	City	Zip Code	Latitude	Longitude
0	0002-ORFBO	Female	37	Yes	0	Frazier Park	93225	34.827662	-118.
1	0003-MKNFE	Male	46	No	0	Glendale	91206	34.162515	-118.
2	0004-TLHLJ	Male	50	No	0	Costa Mesa	92627	33.645672	-117.
3	0011-IGKFF	Male	78	Yes	0	Martinez	94553	38.014457	-122.
4	0013-EXCHZ	Female	75	Yes	0	Camarillo	93010	34.227846	-119.
...	...	...	...	...	...	...	...	...	...
6584	9986-BONCE	Female	36	No	0	Fallbrook	92028	33.362575	-117.
6585	9987-LUTYD	Female	20	No	0	La Mesa	91941	32.759327	-116.
6586	9992-RRAMN	Male	40	Yes	0	Riverbank	95367	37.734971	-120.
6587	9993-LHIEB	Male	21	Yes	0	Solana Beach	92075	33.001813	-117.
6588	9995-HOTOH	Male	36	Yes	0	Sierra City	96125	39.600599	-120.

6589 rows x 38 columns

```
df = df.drop(columns=["Customer ID", "City", "Zip Code","Longitude","Latitude", "Churn Category","Churn Reason"])
```

```
df.isnull().sum()
```

Gender	0
Age	0
Married	0
Number of Dependents	0
Number of Referrals	0
Tenure in Months	0
Offer	0
Phone Service	0
Avg Monthly Long Distance Charges	644
Multiple Lines	644
Internet Service	0
Internet Type	1344
Avg Monthly GB Download	1344
Online Security	1344
Online Backup	1344
Device Protection Plan	1344
Premium Tech Support	1344
Streaming TV	1344
Streaming Movies	1344

```

Streaming Music          1344
Unlimited Data            1344
Contract                  0
Paperless Billing         0
Payment Method           0
Monthly Charge           0
Total Charges            0
Total Refunds            0
Total Extra Data Charges 0
Total Long Distance Charges 0
Total Revenue            0
Customer Status          0
dtype: int64

```

```
df["Avg Monthly Long Distance Charges"] = df["Avg Monthly Long Distance Charges"].fillna(0)
```

```
df["Avg Monthly GB Download"] = df["Avg Monthly GB Download"].fillna(0)
```

```
df["Multiple Lines"] = df["Multiple Lines"].fillna("No Phone")
```

```
df["Internet Type"] = df["Internet Type"].fillna("None")
```

```

na_cols = ["Online Security", "Online Backup", "Device Protection Plan", "Premium Tech Support", "Streaming TV", "Streaming Movies"]
for column in na_cols:
    df[column] = df[column].fillna("No")

```

```
df.isnull().sum()
```

```

Gender          0
Age             0
Married         0
Number of Dependents 0
Number of Referrals 0
Tenure in Months 0
Offer           0
Phone Service   0
Avg Monthly Long Distance Charges 0
Multiple Lines   0
Internet Service 0
Internet Type    0
Avg Monthly GB Download 0
Online Security  0
Online Backup    0
Device Protection Plan 0
Premium Tech Support 0
Streaming TV     0
Streaming Movies 0
Streaming Music  0
Unlimited Data    0
Contract         0
Paperless Billing 0
Payment Method   0
Monthly Charge   0
Total Charges    0
Total Refunds    0
Total Extra Data Charges 0
Total Long Distance Charges 0
Total Revenue    0
Customer Status  0
dtype: int64

```

```
df.shape
```

```
(6589, 31)
```

```
from sklearn.preprocessing import LabelEncoder
```

```

le = LabelEncoder()

for col in df.columns[1:]:
    if df[col].dtype == 'object':
        if len(list(df[col].unique())) <= 2:
            le.fit(df[col])
            df[col] = le.transform(df[col])

```

```
df['Gender'] = [1 if each == 'Female' else 0 for each in df['Gender']]
```

```
def encode_data(dataframe):
    if dataframe.dtype == "object":
        dataframe = LabelEncoder().fit_transform(dataframe)
    return dataframe

data = df.apply(lambda x: encode_data(x))
data.head()
```

	Gender	Age	Married	Number of Dependents	Number of Referrals	Tenure in Months	Offer	Phone Service	Avg Monthly Long Distance Charges
0	1	37	1	0	2	9	0	1	42.36
1	0	46	0	0	0	9	0	1	10.65
2	0	50	0	0	0	4	5	1	33.65
3	0	78	1	0	1	13	4	1	27.86
4	1	75	1	0	3	3	0	1	7.36

5 rows x 31 columns

```
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6589 entries, 0 to 6588
Data columns (total 31 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Gender                                     6589 non-null   int64
1   Age                                       6589 non-null   int64
2   Married                                  6589 non-null   int64
3   Number of Dependents                     6589 non-null   int64
4   Number of Referrals                      6589 non-null   int64
5   Tenure in Months                        6589 non-null   int64
6   Offer                                    6589 non-null   int64
7   Phone Service                           6589 non-null   int64
8   Avg Monthly Long Distance Charges        6589 non-null   float64
9   Multiple Lines                          6589 non-null   int64
10  Internet Service                         6589 non-null   int64
11  Internet Type                           6589 non-null   int64
12  Avg Monthly GB Download                  6589 non-null   float64
13  Online Security                         6589 non-null   int64
14  Online Backup                           6589 non-null   int64
15  Device Protection Plan                  6589 non-null   int64
16  Premium Tech Support                   6589 non-null   int64
17  Streaming TV                           6589 non-null   int64
18  Streaming Movies                       6589 non-null   int64
19  Streaming Music                        6589 non-null   int64
20  Unlimited Data                         6589 non-null   int64
21  Contract                               6589 non-null   int64
22  Paperless Billing                       6589 non-null   int64
23  Payment Method                         6589 non-null   int64
24  Monthly Charge                         6589 non-null   float64
25  Total Charges                          6589 non-null   float64
26  Total Refunds                          6589 non-null   float64
27  Total Extra Data Charges                6589 non-null   int64
28  Total Long Distance Charges             6589 non-null   float64
29  Total Revenue                          6589 non-null   float64
30  Customer Status                        6589 non-null   int64
dtypes: float64(7), int64(24)
memory usage: 1.6 MB
```

```
X = data.drop(['Customer Status'], axis=1)
y = data.loc[:, 'Customer Status'].values
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
X_train.shape

(4612, 30)
```

```
from sklearn.preprocessing import StandardScaler
```

```
cols = ["Tenure in Months", "Avg Monthly Long Distance Charges", "Avg Monthly GB Download", "Monthly Charge", "Total Charges"
```

```
scaler = StandardScaler()  
X_train[cols] = scaler.fit_transform(X_train[cols])  
X_test[cols] = scaler.fit_transform(X_test[cols])
```

```
model = LogisticRegression()
```

```
model.fit(X_train, y_train)
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear\_model/\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge. Increase the number of iterations (max\_iter) or scale the data as shown in:  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(  
 v LogisticRegression  
 LogisticRegression()  
)

```
y_pred = model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)  
precision = precision_score(y_test, y_pred)  
recall = recall_score(y_test, y_pred)  
f1 = f1_score(y_test, y_pred)
```

```
print("Accuracy:", accuracy)  
print("Precision:", precision)  
print("Recall:", recall)  
print("F1-score:", f1)
```

Accuracy: 0.821446636317653  
Precision: 0.8690476190476191  
Recall: 0.8820184790334044  
F1-score: 0.875485008818342