

ML Lab Ex6

1. Use the attached file and run SVM, Decision tree, Random Forest and any one boosting algorithm.
2. Find out the different tunable parameters for each algorithms mentioned above.
3. Apply gridsearchCV and randomizedsearchCV for all the above classification algorithms and get the best parameters.

```
import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn import metrics
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
from scipy.stats import uniform as sp_randFloat
from scipy.stats import randint as sp_randInt
```

```
df = pd.read_csv("/content/Telco-Customer-Churn.csv")
```

df

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneServ
0	7590-VHVEG	Female	0	Yes	No	1	
1	5575-GNVDE	Male	0	No	No	34	
2	3668-QPYBK	Male	0	No	No	2	
3	7795-CFOCW	Male	0	No	No	45	
4	9237-HQITU	Female	0	No	No	2	
...	
7038	6840-RESVB	Male	0	Yes	Yes	24	
7039	2234-XADUH	Female	0	Yes	Yes	72	
7040	4801-JJAZL	Female	0	Yes	Yes	11	
7041	8361-LTMKD	Male	1	Yes	No	4	
7042	3186-AJIEK	Male	0	No	No	66	

7043 rows x 21 columns

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column              Non-Null Count  Dtype
---  -
0   customerID          7043 non-null   object
1   gender              7043 non-null   object
2   SeniorCitizen       7043 non-null   int64
3   Partner             7043 non-null   object
4   Dependents          7043 non-null   object
5   tenure              7043 non-null   int64
6   PhoneService        7043 non-null   object
7   MultipleLines       7043 non-null   object
```

```

8  InternetService  7043 non-null object
9  OnlineSecurity  7043 non-null object
10 OnlineBackup    7043 non-null object
11 DeviceProtection 7043 non-null object
12 TechSupport     7043 non-null object
13 StreamingTV     7043 non-null object
14 StreamingMovies  7043 non-null object
15 Contract        7043 non-null object
16 PaperlessBilling 7043 non-null object
17 PaymentMethod   7043 non-null object
18 MonthlyCharges  7043 non-null float64
19 TotalCharges    7043 non-null object
20 Churn           7043 non-null object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

```
df.isnull().sum()
```

```

customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64

```

✓ 1. Data Preprocessing

```
df = df.drop(["customerID"], axis = 1)
```

```
df["TotalCharges"] = pd.to_numeric(df["TotalCharges"], errors="coerce")
```

```
df.isnull().sum()
```

```

gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    11
Churn           0
dtype: int64

```

```
df = df.fillna(df["TotalCharges"].mean())
```

```
df.isnull().sum()
```

```

gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0

```

```

InternetService      0
OnlineSecurity       0
OnlineBackup         0
DeviceProtection     0
TechSupport          0
StreamingTV          0
StreamingMovies      0
Contract             0
PaperlessBilling     0
PaymentMethod        0
MonthlyCharges       0
TotalCharges         0
Churn                0
dtype: int64

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   gender                7043 non-null   object
 1   SeniorCitizen         7043 non-null   int64
 2   Partner               7043 non-null   object
 3   Dependents            7043 non-null   object
 4   tenure                7043 non-null   int64
 5   PhoneService          7043 non-null   object
 6   MultipleLines         7043 non-null   object
 7   InternetService       7043 non-null   object
 8   OnlineSecurity        7043 non-null   object
 9   OnlineBackup          7043 non-null   object
10   DeviceProtection      7043 non-null   object
11   TechSupport           7043 non-null   object
12   StreamingTV           7043 non-null   object
13   StreamingMovies       7043 non-null   object
14   Contract              7043 non-null   object
15   PaperlessBilling      7043 non-null   object
16   PaymentMethod         7043 non-null   object
17   MonthlyCharges        7043 non-null   float64
18   TotalCharges          7043 non-null   float64
19   Churn                 7043 non-null   object
dtypes: float64(2), int64(2), object(16)
memory usage: 1.1+ MB

```

```
df["PaymentMethod"].unique()
```

```

array(['Electronic check', 'Mailed check', 'Bank transfer (automatic)',
      'Credit card (automatic)'], dtype=object)

```

```
df.replace({'No phone service': "No", 'No internet service': "No"}, inplace=True)
```

```
df["OnlineSecurity"].unique()
```

```
array(['No', 'Yes'], dtype=object)
```

```
df["Churn"].unique()
```

```
array(['No', 'Yes'], dtype=object)
```

```

object_cols = df.select_dtypes(include=['object']).columns
le = LabelEncoder()
for col in object_cols:
    df[col] = le.fit_transform(df[col])
df

```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLi
0	0	0	1	0	1	0	
1	1	0	0	0	34	1	
2	1	0	0	0	2	1	
3	1	0	0	0	45	0	
4	0	0	0	0	2	1	
...
7038	1	0	1	1	24	1	
7039	0	0	1	1	72	1	
7040	0	0	1	1	11	0	
7041	1	1	1	0	4	1	
7042	1	0	0	0	66	1	

7043 rows x 20 columns

Next steps: [Generate code with df](#) [View recommended plots](#)

```
df1 = pd.get_dummies(df)
df1.head()
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLi
0	0	0	1	0	1	0	
1	1	0	0	0	34	1	
2	1	0	0	0	2	1	
3	1	0	0	0	45	0	
4	0	0	0	0	2	1	

Next steps: [Generate code with df1](#) [View recommended plots](#)

```
X = df1.drop(columns = ["Churn"])
y = df1["Churn"].values
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

Model Prediction

SVM

```
svc = SVC(kernel="linear", random_state = 42)
svc.fit(X_train,y_train)
svc_acc = svc.score(X_test,y_test)
print("SVM accuracy:", svc_acc)
```

SVM accuracy: 0.7959431279620853

Decision Tree

```
dt = DecisionTreeClassifier(max_depth = 4)
dt.fit(X_train,y_train)
dt_acc = dt.score(X_test,y_test)
print("Decision Tree accuracy:",dt_acc)
```

Decision Tree accuracy: 0.7856128726928537

Random Forest

```
rf = RandomForestClassifier(n_estimators=100, criterion='gini', random_state=0)
rf.fit(X_train, y_train)
rf_acc = rf.score(X_test, y_test)
print("Random Forest accuracy:", rf_acc)
```

Random Forest accuracy: 0.7931850449597728

```
gb = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0, max_depth=1, random_state=0)
gb.fit(X_train, y_train)
gb_acc = gb.score(X_test, y_test)
print("Gradient Boost accuracy:", gb_acc)
```

Gradient Boost accuracy: 0.808329389493611

✓ GridsearchCV and RandomizedsearchCV

```
params = {'max_depth': [4, 5, 6, 7, 8, 9, 10],
          'criterion': ['gini', 'entropy'],
          'splitter': ['best', 'random'],
          "max_features": [1, 2, 3, 4, 5, 6, 7, 8, 9],
          "min_samples_leaf": [1, 2, 3, 4, 5, 6, 7, 8, 9],
          }
```

✓ RSCV Decision Trees

```
tree = DecisionTreeClassifier(random_state=0)
tree_rscv = RandomizedSearchCV(tree, param, n_jobs=-1)
tree_rscv.fit(X_train, y_train)
print(tree_rscv.best_params_)
```

{'splitter': 'best', 'min_samples_leaf': 6, 'max_features': 8, 'max_depth': 6, 'criterion': 'entropy'}

```
dt = DecisionTreeClassifier(criterion= 'gini', max_depth= 8, max_features= 5, min_samples_leaf= 6, splitter= 'best')
dt.fit(X_train, y_train)
dt.score(X_test, y_test)
```

0.7780407004259347

✓ GSCV Decision Trees

```
tree = DecisionTreeClassifier(random_state=0)
tree_gscv = GridSearchCV(tree, param, cv=10, n_jobs=-1)
tree_gscv.fit(X_train, y_train)
print(tree_gscv.best_params_)
```

{'criterion': 'gini', 'max_depth': 7, 'max_features': 8, 'min_samples_leaf': 8, 'splitter': 'best'}

```
dt = DecisionTreeClassifier(criterion= 'entropy', max_depth= 8, max_features= 9, min_samples_leaf= 8, splitter= 'best')
dt.fit(X_train, y_train)
dt.score(X_test, y_test)
```

0.7818267865593942

✓ Support Vector Machine (SVM)

```
param = {'C': [0.1, 1, 10, 100, 1000],
         'gamma': [1, 0.1, 0.01, 0.001, 0.0001], }
```

✓ RSCV SVM

```
svm = SVC()
svm_rscv = RandomizedSearchCV(svm, param, n_jobs=-1)
svm_rscv.fit(X_train, y_train)
print(svm_rscv.best_params_)
```

{'gamma': 0.0001, 'C': 0.1}

```
svm = SVC( gamma = 0.01, C = 100)
svm.fit(X_train, y_train)
```

```
svm.score(X_test, y_test)
```

```
0.7501183151916706
```

✓ GSCV SVM

```
svm = SVC()
svm_gscv = GridSearchCV(svm, param, n_jobs=-1)
svm_gscv.fit(X_train,y_train)
print(svm_gscv.best_params_)
```

```
{'C': 1, 'gamma': 0.0001}
```

```
svm = SVC( gamma = 0.01, C = 100)
svm.fit(X_train, y_train)
svm.score(X_test, y_test)
```

```
0.7501183151916706
```

✓ Random Forest

```
params = {'n_estimators':[50,75,100,125,200],
          'criterion':['gini','entropy'],
          'bootstrap':[True, False]}
```

✓ RSCV Random Forest

```
rf = RandomForestClassifier(random_state=0)
rf_rscv = RandomizedSearchCV(rf, param, n_jobs=-1)
rf_rscv.fit(X_train,y_train)
print(rf_rscv.best_params_)
```

```
{'n_estimators': 50, 'criterion': 'entropy', 'bootstrap': True}
```

```
rf = RandomForestClassifier(n_estimators=100, criterion='entropy',random_state=0,bootstrap = True)
rf.fit(X_train,y_train)
rfc_accuracy = rf.score(X_test, y_test)
rf.score(X_test, y_test)
```

```
0.7936583057264552
```

✓ GSCV Random Forest

```
rf = RandomForestClassifier(random_state=0)
rf_gscv = GridSearchCV(rf, param, n_jobs=-1)
rf_gscv.fit(X_train,y_train)
print(rf_gscv.best_params_)
```

```
{'bootstrap': True, 'criterion': 'entropy', 'n_estimators': 50}
```

```
rf = RandomForestClassifier(n_estimators=100, criterion='entropy',random_state=0,bootstrap = True)
rf.fit(X_train,y_train)
rfc_accuracy = rf.score(X_test, y_test)
rf.score(X_test, y_test)
```

```
0.7936583057264552
```

✓ Boosting

```
params = {'max_depth'      : [2,3,4,5,6,7],
          'learning_rate': [0.15,0.1,0.05,0.01,0.005,0.001],
          'n_estimators': [100,250,500,750]
        }
```

✓ RSCV Boosting

```
gb = GradientBoostingClassifier(random_state=0)
gb_rscv = RandomizedSearchCV(estimator=gb, param_distributions = param, n_jobs=-1)
gb_rscv.fit(X_train, y_train)
print(gb_rscv.best_params_)
```

```
{'n_estimators': 100, 'max_depth': 2, 'learning_rate': 0.15}
```

```
gb = GradientBoostingClassifier(n_estimators=100, learning_rate=0.05,max_depth=3, random_state=0).fit(X_train, y_train)
gb_accuracy = gb.score(X_test, y_test)
gb.score(X_test, y_test)
```

```
0.7950780880265026
```

▼ GSCV Boosting

```
gb = GradientBoostingClassifier(random_state=0)
gb_gscv = GridSearchCV(gb, param, n_jobs=-1)
gb_gscv.fit(X_train, y_train)
print(gb_gscv.best_params_)
```

```
{'learning_rate': 0.01, 'max_depth': 5, 'n_estimators': 500}}
```

```
gb = GradientBoostingClassifier(n_estimators=500, learning_rate=0.01,max_depth=5, random_state=0).fit(X_train, y_train)
gb_accuracy = gb.score(X_test, y_test)
gb.score(X_test, y_test)
```

```
0.7980679289099526
```